

inference

Quantization

→ Caching

→ Speculative Decoding

→ Continuous batching

→ Pruned Attention

with C & S The time cost of inference is dominated

memory
latency
throughput



role: 'state'

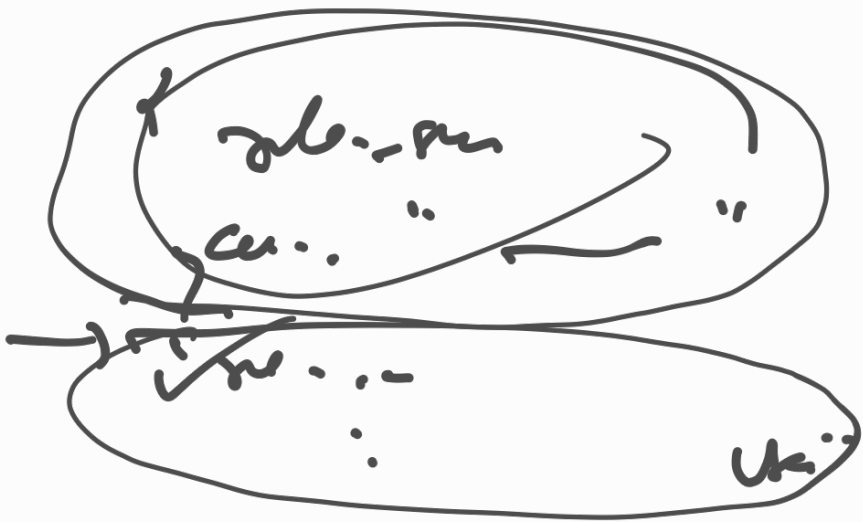
symbol: 'quax helix count'

role: 'inr'

cost: 'used in the context table'



Amended. The center of the hole is the



Ca clip

three
land

even
attacker matches
with a - 20%

pre-~~filter~~
point level

even
attacker
matches



pre-
filter
+
decal

even

4's

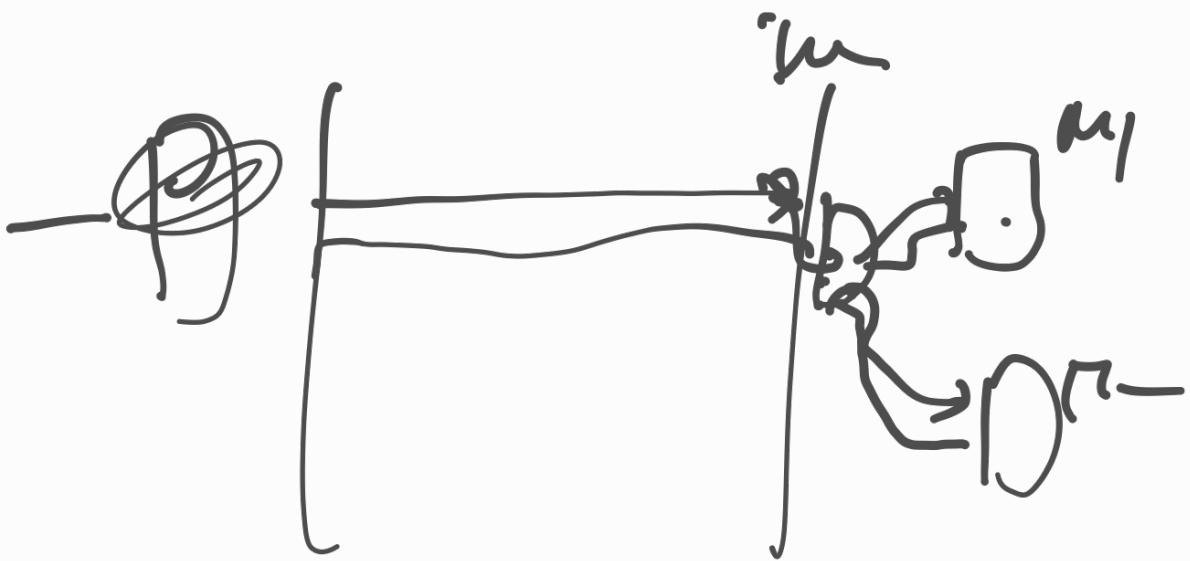
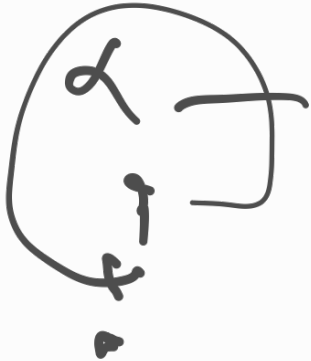


the


as
a/cd
as/cd
am



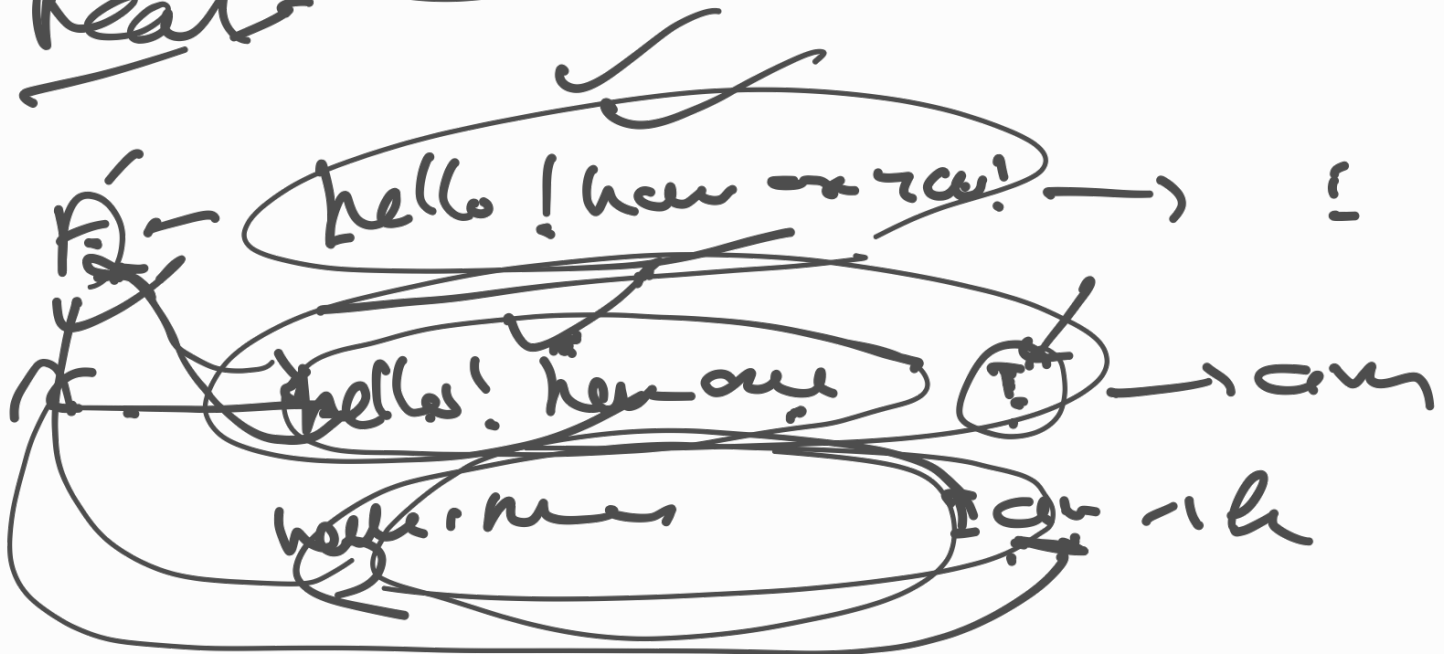
Reader 22



def func(i, v, (m, s))



near



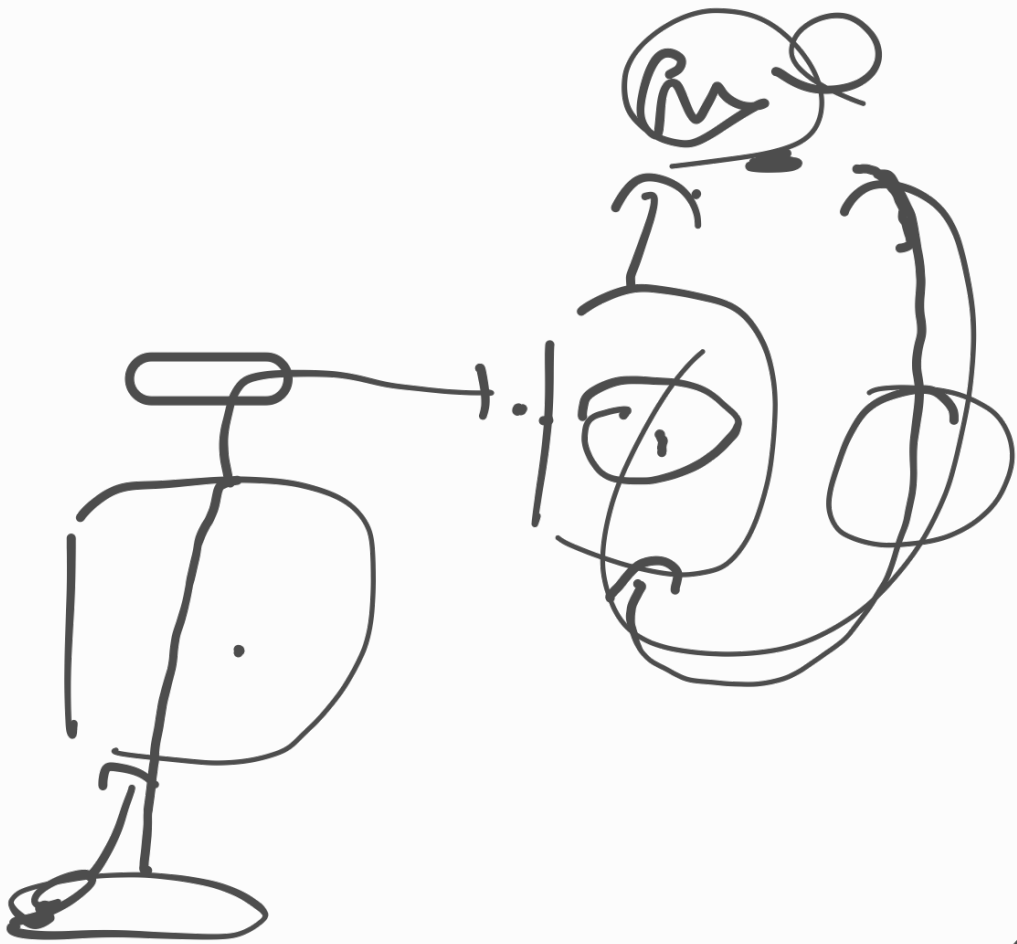
hey

(2)

encoder



encoder decoder



SLG

Control
Control system



Speculative Decale

